

Automated Recognition of Wasp Species through Image Analysis

Adrian Teng-Amnuay
Taiwan Forestry Research Institute
Taipei, Taiwan

Presentation Layout

- Introduction
 - Background
 - Project description and motivation
 - Designing a solution
 - Related work and available tools
 - Describe the workflow

Presentation Layout

- System implementation
 - (Manual) Image Preprocessing
 - (Automatic) Image preprocessing
 - Feature Extraction
 - Color based features
 - Shape based features
 - Feature reduction
 - Machine learning
 - Model validation
 - Choices for statistical models
 - Model evaluation

Presentation Layout

- Conclusion
 - Results Analysis
 - Unaddressed concerns
 - Future work
 - Discussion

Background

- Adrian Teng-Amnuay
- Undergraduate student from UC San Diego
 - Pursuing B.S. in Computer Engineering (June 2012)
- Before this project
 - One course in Computer Vision
 - One course in Artificial Intelligence
 - Studied Machine Learning with Tony Fountain

Defining the Project

- Image classification of Vespidae Wasps
- Given a set of images, can we determine what species of wasp is in each image?



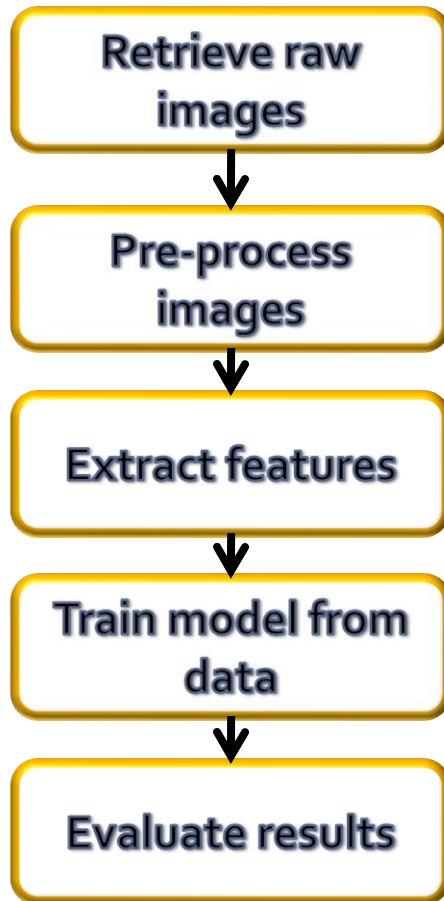
Motivation

- Why do we want an automated image classification system?

How to accomplish this goal?

- Study of Computer Vision (CV) “eyes” and Machine Learning (ML) “brain”
- Can a machine learn to recognize patterns in image data?
- Classification is just one application of CV/ML
- Many CV and ML algorithms already implemented in the OpenCV library¹

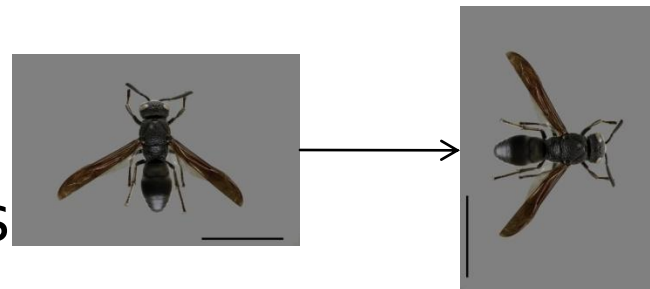
Image classification workflow



- Preprocessing
 - “Garbage in, garbage out”
 - Remove as much irrelevant data as possible
- Feature Extraction
 - How to “describe” an image?
 - Our system focuses on color and shape
- Statistical models of learning
 - How to learn and compare features?
 - How to evaluate performance?
 - Does model generalize to real-world data?

Manual image preprocessing

- Why not automate?
 - Either too hard to automate or easier to do manually
- Image compression
 - Hundreds of high-resolution images (~2 MB each)
- Rotating images
 - For consistency of images



Manual image preprocessing

- Segment out wings, legs, antennae, and knife-blade
 - Too much intraclass variance
 - Difficult to automatically detect and segment



Wings roughly 45°



Wings closed, body obscured.



Wings wide open



Knife blade protruding

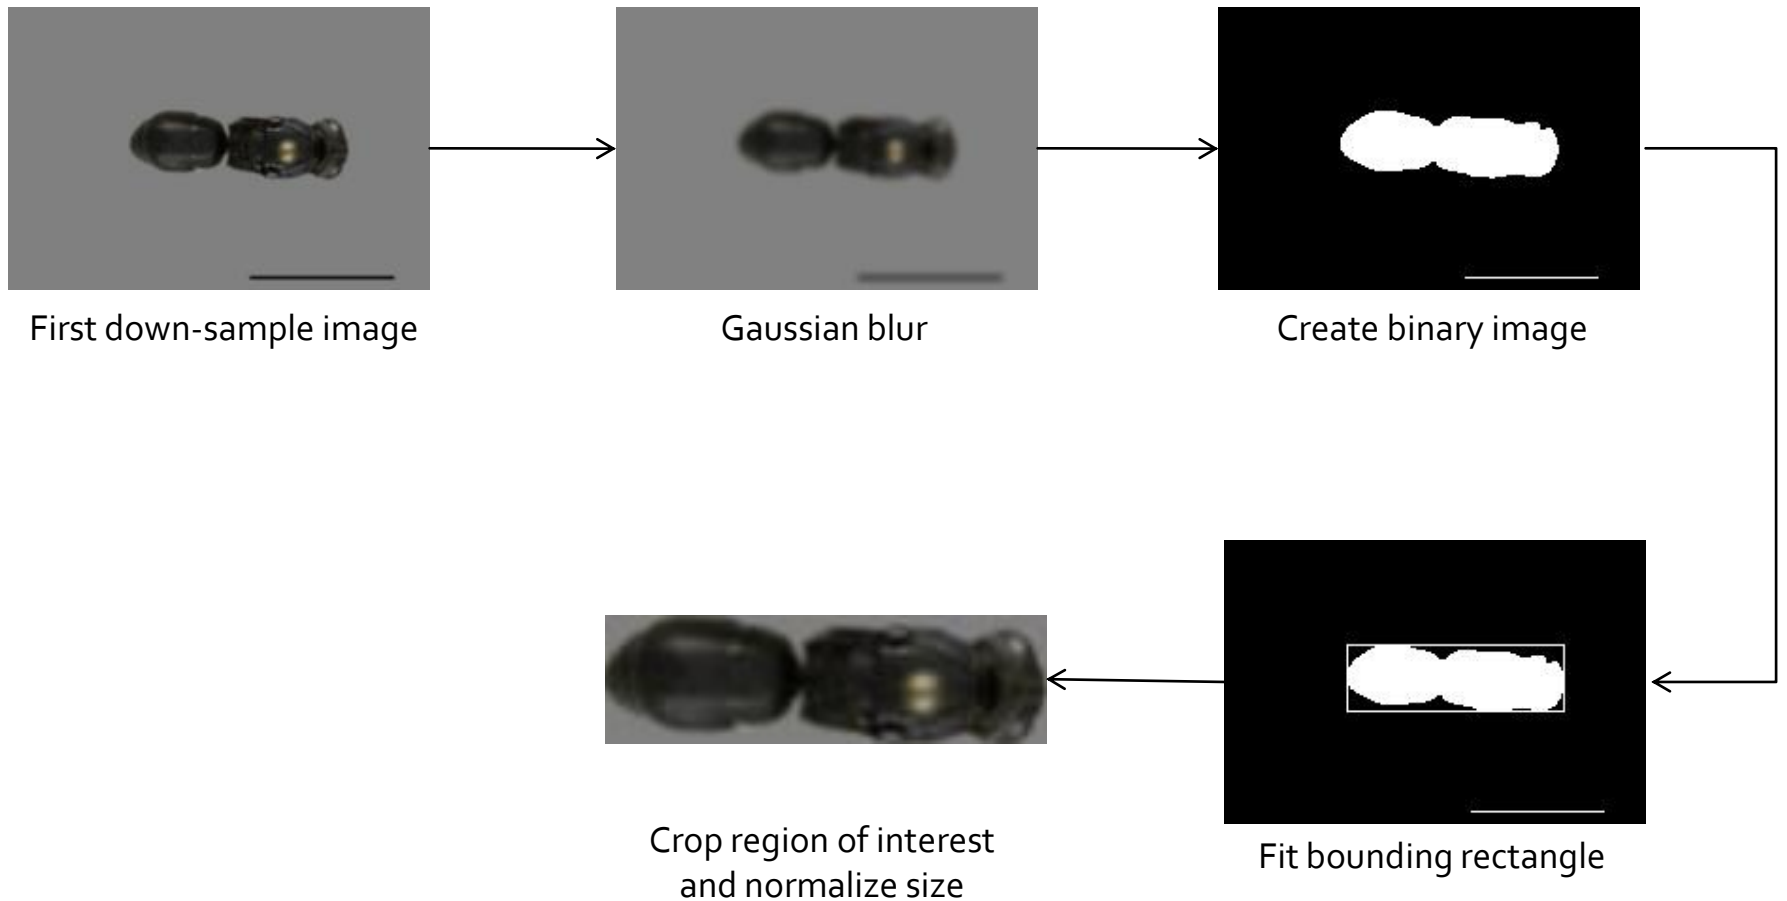
Manual image preprocessing

- Preserved wasp bodies



- Use Photoshop to remove unwanted parts

Automatic image preprocessing



Feature Representation (example)

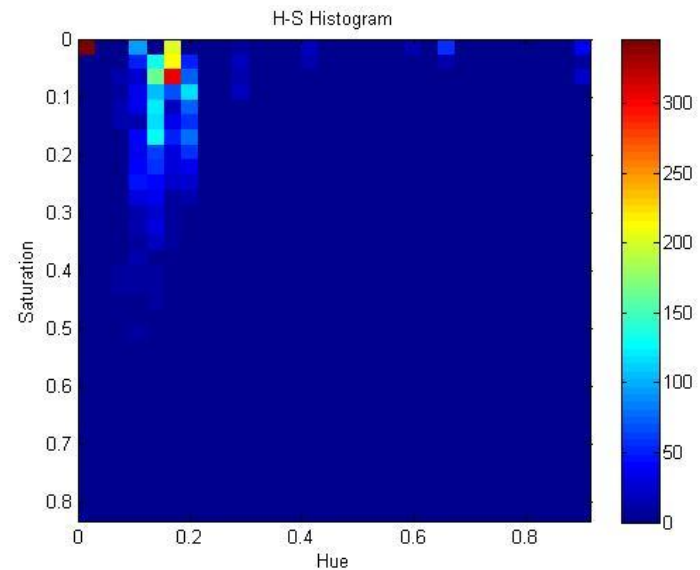
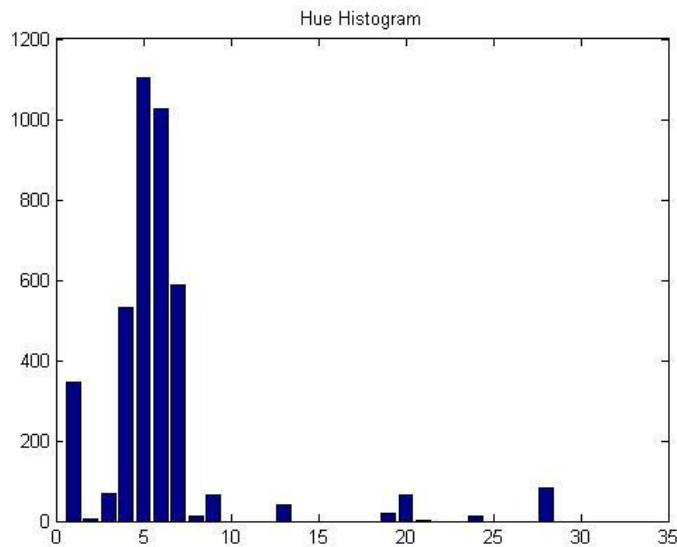
- Weight
- Length

Feature Extraction

- How to describe color?
- Color Histograms
 - Compute 2D Hue-Saturation histograms
 - Disregard “Value” to remove sensitivity to lighting

Color Histograms

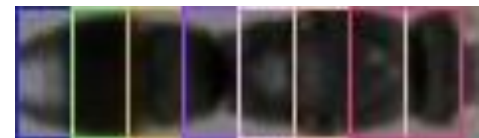
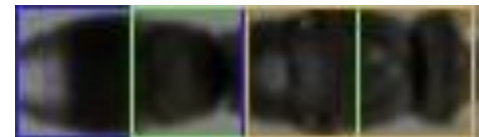
- Examples 1D and 2D histograms



Together, hue and Saturation
provide a good “description”
of color

Overlapping blocks of Color Histograms

- Color histograms place pixels from the entire image into bins
 - No spatial information
- To capture spatial information about color, we take blocks from different regions of the image and calculate color histograms for each block.
- We can vary the block sizes, and allow the blocks to overlap to improve performance.



Feature Extraction

- How to describe shape?
- Histogram of Oriented Gradients (HOG)²
 - Robust shape descriptors created for pedestrian detection in images
 - Already implemented by OpenCV

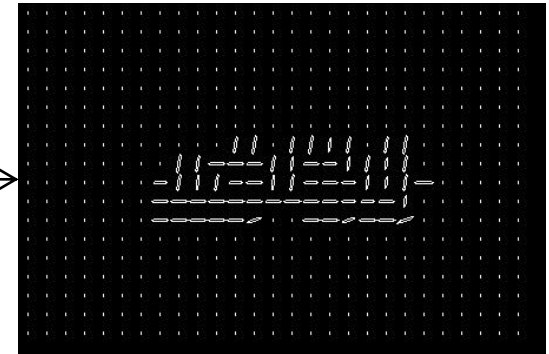
² N. Dalal and B. Triggs. Histograms of oriented gradients for human detection.

Histogram of Oriented Gradients



For the sake brevity, algorithm details are not covered

HOG features are explained in detail in the paper by N. Dalal and B. Triggs [2]



Cells of grouped pixels, represented by magnitude and orientation

² N. Dalal and B. Triggs. Histograms of oriented gradients for human detection. In *CVPR*, pages 1: 886–893, 2005.

Feature Reduction

- Color and HOG feature vectors contain thousands of features (over 10,000 with the current implementation)
- Curse of dimensionality
 - Feature vectors become difficult to compare
 - Computational overhead grows exponentially with number of features
- Principle Component Analysis (PCA)
 - A way to reduce feature dimensionality
 - Principle components only need to be calculated during training (takes ~30 seconds on Intel® Core™2 Duo Processor T5550 (1.83 GHz))

Choosing a model

- Many different ML models
 - Nearest Neighbor
 - K-nearest neighbor
 - Decision Trees
 - Support Vector Machines
 - Binary classifier
 - Multiclass SVM (one-against-one, one-against-rest)
 - Many more, but only had time to test out a handful of models

Model Validation

- Need a way to evaluate model performance
- Holdout validation
- K fold cross-validation (10)
- Confusion Matrices
 - How often is a class correctly predicted?
 - How often is a class confused for another class?

Evaluating the model

- One-against-rest SVMs using overlapping blocks of color histograms
- Validation accuracy in percentage \pm standard deviation across 10 folds
- (number correctly identified / total number of images)

	0	1	2	3	4	5	6	7	8	9	10	11	12	13	14	15	# of Test Images	Accuracy
0	13	0	0	0	0	0	0	0	0	0	0	0	0	0	0	0	13	1
1	0	12	0	1	0	0	0	0	0	0	1	0	0	0	0	0	14	0.86
2	1	0	38	0	0	0	0	1	0	0	2	0	0	1	0	0	43	0.88
3	0	0	0	47	0	0	0	0	0	0	0	0	0	0	0	1	48	0.98
4	0	0	0	0	40	0	0	0	0	0	0	0	0	0	0	0	40	1
5	0	0	0	0	0	18	1	0	0	0	0	0	0	0	0	0	19	0.95
6	0	0	0	0	0	1	24	0	0	0	0	0	0	0	0	0	25	0.96
7	0	0	2	1	0	0	0	16	0	0	1	0	0	1	0	0	21	0.76
8	0	0	0	1	0	1	0	0	9	0	0	0	0	0	0	0	11	0.82
9	0	0	0	0	0	0	0	0	0	21	0	0	0	0	0	0	21	1
10	0	0	1	0	0	0	0	0	0	0	7	0	0	1	0	1	10	0.7
11	0	0	0	0	0	0	0	0	0	0	0	33	1	0	0	0	34	0.97
12	0	0	0	0	0	0	0	0	1	0	0	3	31	0	0	0	35	0.89
13	0	0	2	0	0	0	0	0	0	0	0	0	0	44	0	0	46	0.96
14	0	0	0	1	0	0	0	0	0	0	0	0	0	0	42	1	44	0.95
15	0	0	0	1	0	0	0	0	0	0	0	0	0	0	1	11	13	0.85

Color features:
 $92.906 \pm 4.534 \%$
 (406/437)

	0	1	2	3	4	5	6	7	8	9	10	11	12	13	14	15	# of Test Images	Accuracy
0	10	0	0	0	0	0	0	1	0	0	0	0	0	2	0	0	13	0.77
1	0	12	0	0	0	0	0	0	0	0	0	0	0	0	1	1	14	0.86
2	0	0	38	0	0	0	0	4	0	0	1	0	0	0	0	0	43	0.88
3	0	0	2	44	0	0	0	1	0	1	0	0	0	0	0	0	48	0.92
4	0	0	0	0	37	0	0	0	0	0	1	0	2	0	0	0	40	0.93
5	0	0	0	0	0	17	2	0	0	0	0	0	0	0	0	0	19	0.89
6	0	0	0	0	2	3	19	0	1	0	0	0	0	0	0	0	25	0.76
7	0	0	5	0	0	0	0	13	0	0	1	0	0	1	1	0	21	0.62
8	0	0	1	0	0	1	1	0	6	0	1	1	0	0	0	0	11	0.55
9	0	0	0	1	0	0	1	0	0	19	0	0	0	0	0	0	21	0.9
10	1	1	0	0	0	0	0	1	0	0	7	0	0	0	0	0	10	0.7
11	0	0	0	0	0	0	0	0	0	0	0	33	1	0	0	0	34	0.97
12	0	0	0	0	2	1	0	0	0	0	0	1	31	0	0	0	35	0.89
13	1	0	3	0	0	0	0	1	0	0	0	0	0	41	0	0	46	0.89
14	0	0	0	0	0	0	1	0	0	1	0	0	1	0	41	0	44	0.93
15	0	0	0	0	0	0	0	0	0	0	0	0	0	1	2	10	13	0.77

HOG features:
 $86.499 \pm 4.843 \%$
 (378/437)

Combining color and shape features

- Instead of only predicting the class label for a given example, return probabilistic output for each class
 - e.g. “an_image.jpg” is 98% likely to be “Allorhynchium argentatum,” 0.1% likely to be “Apodynerus f.formosensis,” 0.3% likely to be ... and so on.
- This allows us to combine shape and color features in a naïve fashion

Evaluating the model

- One-against-rest SVMs using both color blocks and HOG features
- Validation accuracy in percentage \pm standard deviation across 10 folds
- (number correctly identified / total number of images)

	0	1	2	3	4	5	6	7	8	9	10	11	12	13	14	15	# of Test Images	Accuracy
0	13	0	0	0	0	0	0	0	0	0	0	0	0	0	0	0	13	1
1	0	14	0	0	0	0	0	0	0	0	0	0	0	0	0	0	14	1
2	0	0	41	0	0	0	0	1	0	0	1	0	0	0	0	0	43	0.95
3	0	0	0	47	0	0	0	0	0	0	0	0	0	0	0	1	48	0.98
4	0	0	0	0	40	0	0	0	0	0	0	0	0	0	0	0	40	1
5	0	0	0	0	0	18	1	0	0	0	0	0	0	0	0	0	19	0.95
6	0	0	0	0	0	2	23	0	0	0	0	0	0	0	0	0	25	0.92
7	0	0	2	1	0	0	0	17	0	0	1	0	0	0	0	0	21	0.81
8	0	0	1	0	0	0	0	0	10	0	0	0	0	0	0	0	11	0.91
9	0	0	0	0	0	0	0	0	0	21	0	0	0	0	0	0	21	1
10	0	0	1	0	0	0	0	0	0	0	9	0	0	0	0	0	10	0.9
11	0	0	0	0	0	0	0	0	0	0	0	34	0	0	0	0	34	1
12	0	0	0	0	0	0	0	0	0	0	0	1	34	0	0	0	35	0.97
13	0	0	0	0	0	0	0	0	0	0	0	0	0	46	0	0	46	1
14	0	0	0	0	0	0	0	0	0	0	0	0	1	0	42	1	44	0.95
15	0	0	0	1	0	0	0	0	0	0	0	0	0	0	1	11	13	0.85

96.110 \pm 1.810 %
(420/437)

Results Analysis

- We have presented a workflow for automated image classification, and achieved an overall performance of ~96%
- Using local color information and HOG descriptors for shape captures much of the interclass variance
- Still room for improvement

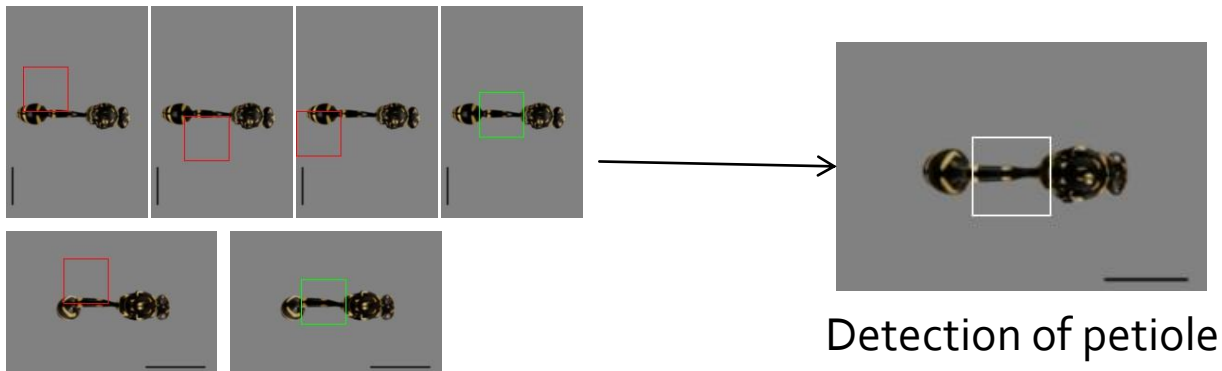
Unaddressed concerns

- Tuning parameters
 - Large number of parameters tested and chosen manually, not optimally
 - Number of bins for color histograms
 - HOG parameters
 - Cropped image resizing dimensions
 - SVM parameters
 - Probability output function parameters
- Number of features
 - Shouldn't have 10,000 features
- Scale bar too close to wasp body
 - Leads to bad binary image



Future work

- Automating the entire workflow
- An important step that was simplified in the problem was manual segmentation of wings, legs, etc.
 - To automate, look into part-based detection as described in Week 6 Report



- But train on entire wasp body, focusing on wings as negatives
- Feed detected "wasp body" into current system

Future work

- Modify current implementation to work with wings
 - Might not be necessary to remove wings
 - Other creative ideas? Tweak system implementation?
- Expand system to include full set of species
 - Did not include several species due to small number of training samples
- Add graphical user interface
 - Include visual display of results (confusion matrices, species labels, etc.)

Discussion

- Any questions?

References

- ¹ G. Bradski. The OpenCV Library. *Dr. Dobb's Journal of Software Tools*, 2000
- ² N. Dalal and B. Triggs. Histograms of oriented gradients for human detection. In *CVPR, pages 1: 886–893, 2005*.

Acknowledgements

- *Mentors*
 - Tony Fountain (CALIT²)
 - Chau Chin Lin (TFRI)
 - Sheng-Shan Lu (TFRI)
 - Yu-Huang Wang (TFRI)
- *PRIME coordinators*
 - Gabriele Wienhausen
 - Teri Simas
 - Peter Arzberger
 - Tricia Taylor
- *Advice on Computer Vision*
 - Serge Belongie (UCSD faculty)

